

# Sentiment Analysis Stock Trading Via Twitter

Jason Yao, Jenna Denker, James Zhang, Yu Li

Department of Computer Science,  
Courant Institute of Mathematical Sciences

New York University,

New York, NY, USA

{JasonYao, jld508, yjz208, yl2133}@nyu.edu

May 5, 2017

## Abstract

The utilization of Sentiment Analysis (SA) on twitter tweets in our system has generated profits of 1.8% over a three week period in April 2017. The market growth during that same time period of the S&P index was 0.02%. This performance increase was derived from leveraging public sentiment about public companies, which does factor into the stock price of public companies. This paper will explore the specific algorithm utilized to derive these market gains, as well as future potential improvements to improve profit yields over time.

## 1 Introduction

We use a custom built program built in Python 3, leveraging Twitter's API in order to create an automated trade file, which contains a list of trades to be executed. Since the building of a whole trading engine is wildly out of scope for this analysis, we thus manually enter the generated trades into a paper trading platform to show the potential that SA has in regards to company stock performance.

## 2 Stocks Market Domain Knowledge

### 2.1 Call Options

A **call option** is an agreement between two investors that gives one the option, but not the obligation, to buy a stock of a company at a set price, known as the strike price within a certain time-frame. (Securities, 2017a) Normally this occurs when an investor believes that a company's stock price will increase in the future, allowing the investor the ability to buy the stock at the current lower price, then sell the stock if the company's stock price has increased.

If, however, a company's stock value has decreased during this time, then there is no reason for the investor to buy at the strike price, when buying on the open market is cheaper. In this case, the investor will have lost the amount paid as a premium for the right to buy the stock at the strike price.

For example, if an investor bought a call option for 500 shares of IBM Stock with a strike price at \$100 when the share price is at \$90, set to expire in 2 months, the investor would theoretically pay \$5000 as their premium cost ( $\$10 \times 500$  shares). If the stock price increased to \$100, the investor will still be at a loss for \$5000 due to the premium cost. The investor will begin to make a profit when the IBM stock price exceeds \$110, when the \$10 gross profit covers the \$5000 premium cost.




Figure 1: Example of a Call Option

### 2.2 Put Options

A **put option** is an agreement between two investors that gives one the option, but not the obligation, to sell a stock of a company at a set price, known as the strike price within a certain time-frame. (Securities, 2017b) Normally this occurs when an investor believes that a company's stock price will decrease in the future, allowing the in-

vestor the ability to sell the stock at the current higher price, and can be thought of as the opposite of a call option.

If, however, a company's stock value has increased during this time, then there is no reason for the investor to sell at the strike price. In this case, the investor will have lost the amount paid as a premium for the right to sell the stock at the strike price.

For example, if an investor bought a call option for 500 shares of IBM Stock with a strike price at \$100 when the share price is at \$90, set to expire in 2 months, the investor would theoretically pay \$5000 as their premium cost (\$10 x 500 shares). If the stock price increased to \$100, the investor will still be at a loss for \$5000 due to the premium cost. The investor will begin to make a profit when the IBM stock price exceeds \$110, when the \$10 gross profit covers the \$5000 premium cost.

For example, if an investor is bearish on UAL, which has a share price of \$75, the investor may put a put option for a period of 1 month for \$70 per share, plus a premium cost of \$1 per share. This means that in total, the investor has bought 500 shares and paid \$500 as a premium cost. If the stock price fell to \$70, then the investor would have still lost the \$500 premium cost. It's only if the stock price falls under to \$69, from which the investor begins to make a profit. On the other hand, if the stock price rose beyond \$70, the investor would let the option expire unexercised, and thus limiting the total cost to the investor to just the initial premium cost of \$500.




Figure 2: Example of a Put Option

## 2.3 Strike Price

The **strike price** for an option is the fixed amount at which an investor agrees to either buy or sell

an underlying asset such as a stock option. For call options, the strike price is where the security can be bought (up to the expiration date); for put options, the strike price is the price at which shares can be sold. (Picardo, 2017)




Figure 3: Example of a Strike Price

## 2.4 Algorithmic Strike Price Determination

When manually entering trades into the paper trading engine, the strike price was automatically determined based off of the existing price, and a simple percentage-based offset.

We divide each option into two categories: "Soft" and "Hard". Soft options (soft calls and soft puts) are generated in the trade file whenever the general sentiment around a company affects its **Viability Score (VS)** to the point of passing the soft threshold, which in our program was set to  $\pm 200$  VS. Likewise, hard options (hard calls and hard puts) are generated in the trade file whenever the general sentiment around a company affects its VS to the point of passing the hard threshold, which was set to  $\pm 500$  VS.

What this means is that if the general sentiment surrounding a company becomes negative extremely quickly, then the company's corresponding VS will drop extremely quickly as well, allowing the program to deduce from the general negative sentiment that a negative public event has occurred (e.g. A PR fiasco, or a bad earnings call), and to react to this event by generating a hard put.

On the other hand, assuming that there's a general positive sentiment trend for a company, its VS score will gradually increase enough to the point where a soft call will be generated instead.

## 2.5 Stock Sentiment Bias

In the stock market, the number of investors who are willing to put money into stocks with positive sentiment and outlook is much higher than people who are willing to play on bad news. This is likely due to the fact that the maximum profit from a short is simply 100% of a company's stock value, which occurs extremely rarely.

On the other hand, however, there is no profit limit for having long call positions, due to the potential growth in value of the stock being unlimited in theory. Thus it should be noted that a wave of positive sentiment would likely drive the stock price of a company up, while negative sentiment would not likely drive the stock down as much. This means that the overpricing of a stock's value is much more likely than its underpricing in the market.

## 3 System Overview

The current system build is segregated into two portions: the Ingestion Engine and the Processing Engine. Due to it being extremely out of scope for this project, a Trading Engine was not built, though one in theory could create one in order to take the trade file generated by the Processing Engine in order to automatically place trades on a market, thus fully automating the system.

### 3.1 Ingestion Engine

The Ingestion Engine is a simple command-line wrapper that pulls Tweets from Twitter via its Streaming API (Twitter, 2017). Every 100 tweets are then batched together for processing in the Processing Engine, and future batches are queued up as required.

Due to the large amount of tweets being ingested into the system itself, the general amount of time required to process each batch is close to 10 seconds. A way to mitigate this required processing time would be to load-balance incoming tweets, and spin up processing servers as required to process and update company VS scores.

### 3.2 Processing Engine

The Processing Engine parses the given input (tweets), combines it with historical data, then computes and labels each given stock with a Viability Score (VS). This score is calculated from the given current sentiment about a stock, and its

current position in the marketplace, with a higher differential resulting in a higher VS.

## 4 Performance

This algorithmic approach to stock trading based off of Twitter users' sentiment analysis derived gains of 1.8%. In this same time period, the S&P index gained only a 0.02% increase, showing that it is at least possible to generate a profit based off of the sentiment analysis of Twitter users that can beat the market in the short term.

In general, because of the deluge of tweets coming into the system in real time, it is not possible to process the tweets in a fast enough manner using a single machine. However, if the processing task was load-balanced with multiple machines handling the processing of information, it is then theoretically possible to analyze the incoming tweets in real-time, instead of having to batch and process tweets in the current system.

### 4.1 Viability Score

The **Viability Score** (VS) is a term used to denote the "sentiment value" of a company over time. As tweets stream into the system, the VS of a given company will increase or decrease to correspond with the general sentiment around a company. When this sentiment reaches a certain threshold, a trade will be generated, and written into the trade file. If a hard threshold is passed, the VS will reset back to 0 to allow for further trades to be conducted.

## 5 Algorithm

For each tweet, we do a first pass to extract out any company names that can be found in the tweet, based off of a pre-defined list of stock tickers, full company names, and a selection of colloquial names of companies.

After the first pass, assuming that at least one company name was found, we then leverage the Python library Textblob in order to gain the sentiment analysis polarity and certainty of the tweet.

The SA polarity is a floating point value between -1 and 1, such that -1 corresponds to an extremely low sentiment, while a value of 1 corresponds to an extremely high sentiment. An SA polarity score of 0 is neutral, and thus does not truly affect the company in question.

The SA certainty is a floating point value between 0 and 1, such that 0 corresponds to an ex-

tremely uncertain analysis, and 1 corresponds to an extremely certain analysis. This is useful in determining whether a tweet is truly as positive or negative as its polarity value would have us believe. High certainty scores thus affects the VS generation much more than low certainty scores.

After the SA polarity and certainty have been extracted, the **Viability Score** (VS) is then generated from by the algorithm below:

$$VS = VS_{old} + (SA_{polarity} * SA_{certainty} * \ln(impressions))$$

## 6 Future Improvements

First and foremost, the biggest improvement to the current system is the implementation of an automated trading engine. In theory, the automated trading engine would allow the system to run end-to-end, taking the companies with high VS, and executing the trades immediately in order to lessen the delay when it comes to having a human investor manually execute the given trades in the trade file.

Another potential improvement to further expand the existing Ingestion Engine would be to implement monitoring of trending hashtags on Twitter. Trending hashtags would allow the system to more accurately monitor breaking news and extreme cases which might not be caught as fast by the user list currently implemented. These extreme cases could include public relations disasters, viral marketing campaigns, and even terrorist attacks, each of which tend to have a large impact on the state of the stock market, and the companies in them. Tweepy has existing functionality to retrieve trending hashtags, and the hashtags could be extracted alongside each batch of tweets and processed accordingly.

Finally, another potential improvement to further expand the capabilities of the system is to include the use of a neural network in order to fine-tune thresholds to be company-specific. Currently, all thresholds are globally applicable to all companies, but it's easy to see situations in which this is not acceptable. Highly popular companies such as Tesla or Facebook are talked about much more frequently than others, and having rarely-talked about companies explode in sentiment value should generate a higher VS, and have lower thresholds due to their rare use. This would also have the benefit of lessening false positives, and would allow the system to better understand

whether a trade is a good choice based off of historical, financial, and sentiment data, along with any extruding market factor at the time of the trade.

## 7 Conclusion

By analyzing the sentiment around Twitter users, it is possible to generate a "sentiment value" score about publicly traded companies, and to then use these trades in order to generate a profit that could beat the market in the short term. The current paper profit yield of 1.8% while the S&P index had a growth of 0.02% showcases the ability to leverage sentiment analysis in the pursuit of intra-day profits.

## 8 References

### References

Elvis Picardo. 2017. Options basics: How to pick the right strike price.

Kotak Securities. 2017a. Understanding how call options work.

Kotak Securities. 2017b. Understanding how put options work.

Twitter. 2017. Twitter streaming api.